

kasan 介绍

Tong.Chen

内容

- 1.功能
- 2.原理及实现
- 3.patch 整合
- 4.使用
- 5.参考

功能

✦ 运行时的内存访问有效性检查

1. Use-after-freed 错误(uaf)

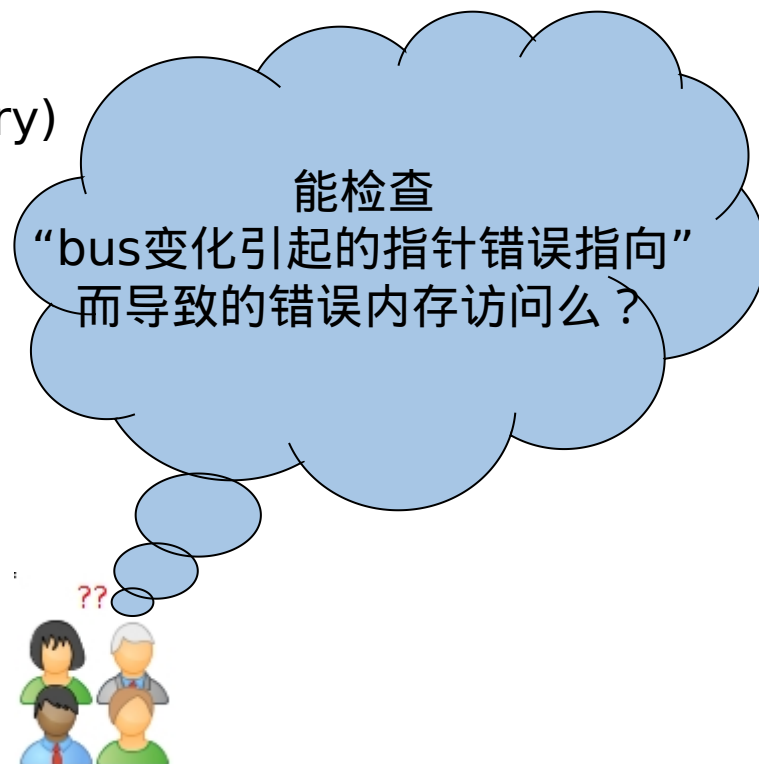
- a. slab 对象
- b. struct page(low memory)
- c. mempool 元素

2. Out-of-bounds 错误(oob)

- a. 1中所有
- b. stack 上的局部变量
- c. 全局数据(!)

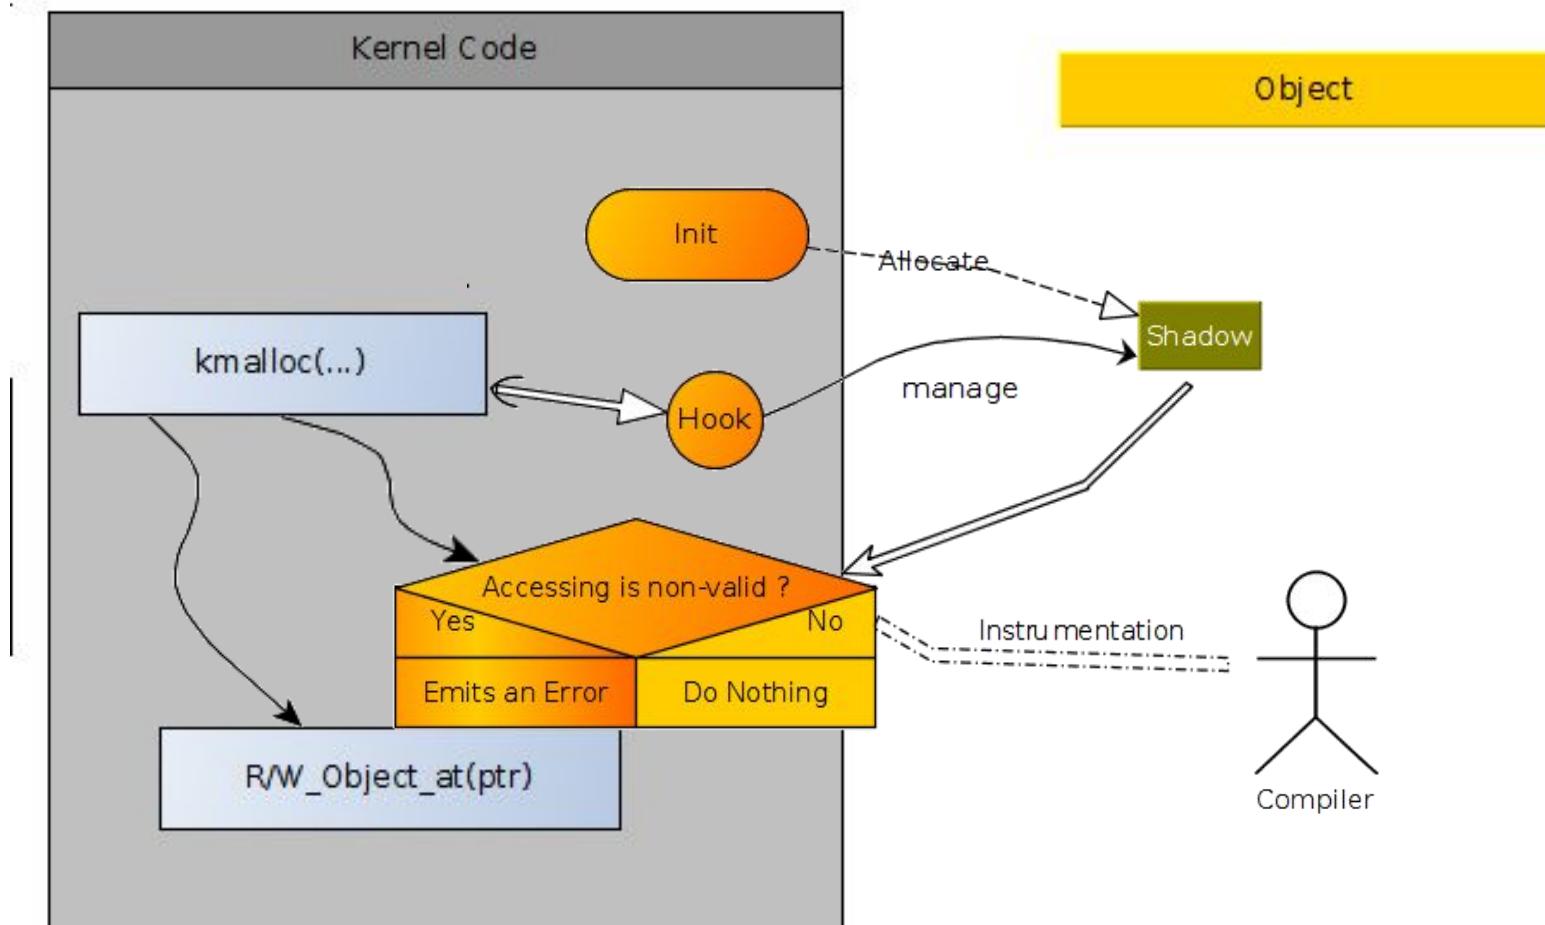
✦ 当前不支持的功能

- 1. vmalloc (module)
- 2. 未初始化内存的访问




原理及实现

✦ 原理



原理及实现



✦ 实现 (ARM 32bits)

1. 架构相关部分

a. 初始化保留 512M

b. 分两阶段进行

- ① 首先暂时得先映射所有 shadow mem 至一个 zero page
- ② 等 low mem 地址空间准备好, 再映射至实际内存

```
60 static inline const void *kasan_shadow_to_mem(const void *shadow_addr)
61 {
62     return (void *)(((unsigned long)shadow_addr - KASAN_SHADOW_OFFSET)
63                     << KASAN_SHADOW_SCALE_SHIFT);
64 }
```

c. Shadow 字节格式

- ① 正数 $n([1,7])$ 表示能访问前面 n 个字节
- ② 0 表示所有 8 字节都能访问
- ③ 负数表示所有 8 字节都不能访问

原理及实现

✦ 实现 (ARM 32bits) (续1)

2. 实现的 Hooks

a. for slub

kasan_slab_alloc / kasan_slab_free
kasan_kmalloc_large / kasan_kfree_large

b. for buddy

kasan_alloc_pages / kasan_free_pages

c. for mempool

kasan_krealloc / kasan_kfree

```
static void kasan_poison_shadow(const void *address, size_t size, u8 value) {...}
void kasan_unpoison_shadow(const void *address, size_t size)
{
    kasan_poison_shadow(address, size, 0);

    if (size & KASAN_SHADOW_MASK) {
        u8 *shadow = (u8 *)kasan_mem_to_shadow(address + size);
        *shadow = size & KASAN_SHADOW_MASK;
    }
}
```

原理及实现

✦ 实现 (ARM 32bits) (续2)

3. Instrumentation

a. 实现一系列函数直接给 compiler 自动做 Instr

- ① __asan_load1 / __asan_store1
- ② __asan_load2 / __asan_store2
- ③ __asan_load4 / __asan_store4
- ④ __asan_load8 / __asan_store8
- ⑤ __asan_load16 / __asan_store16
- ⑥ __asan_loadN / __asan_storeN

思考：
为什么这里有 N() 的版本？

```
static __always_inline void check_memory_region(unsigned long addr,
                                                size_t size, bool write)
{
    struct kasan_access_info info;
    unsigned int kf;

    //...
    if (likely(!memory_is_poisoned(addr, size)))
    {
        return;
    }

    kasan_report(addr, size, write, _RET_IP_);
}
```

原理及实现

✦ 实现 (ARM 32bits) (续3)

3. Instrumentation(续)

a. 手动地为memset/memmove/memcpy 做 Instr

① arch-dependent, asm, non_instrable

② weaken original, use our instred version instead

```
.weak memcpy
ENTRY(__memcpy)
ENTRY(memcpy)

#include "copy_template.S"
```

```
ENDPROC(memcpy)
ENDPROC(__memcpy)

#undef memcpy
void *memcpy(void *dest, const void *src, size_t len)
{
    __asan_loadN((unsigned long)src, len);
    __asan_storeN((unsigned long)dest, len);

    return __memcpy(dest, src, len);
}
```


原理及实现

✦ 实现 (ARM 32bits) (续4)

4. 对 Low Mem Page 的check支持

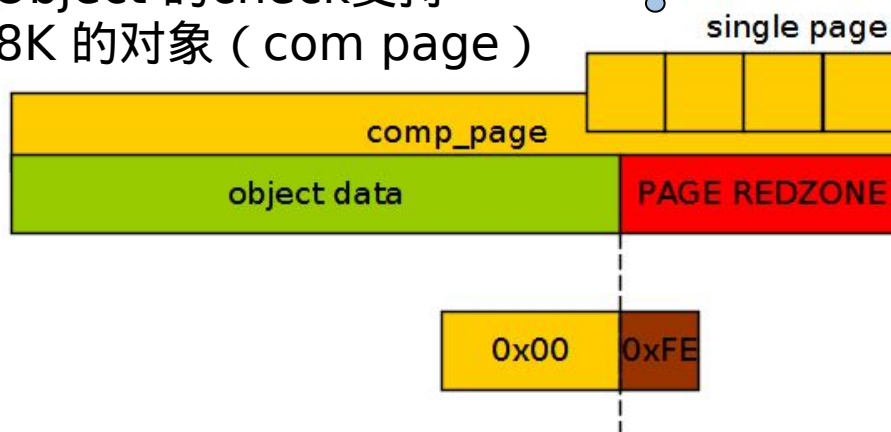


思考：
这里可能
存在漏检么？



5. 对 SLUB Object 的check支持

a. 大于 8K 的对象 (com page)

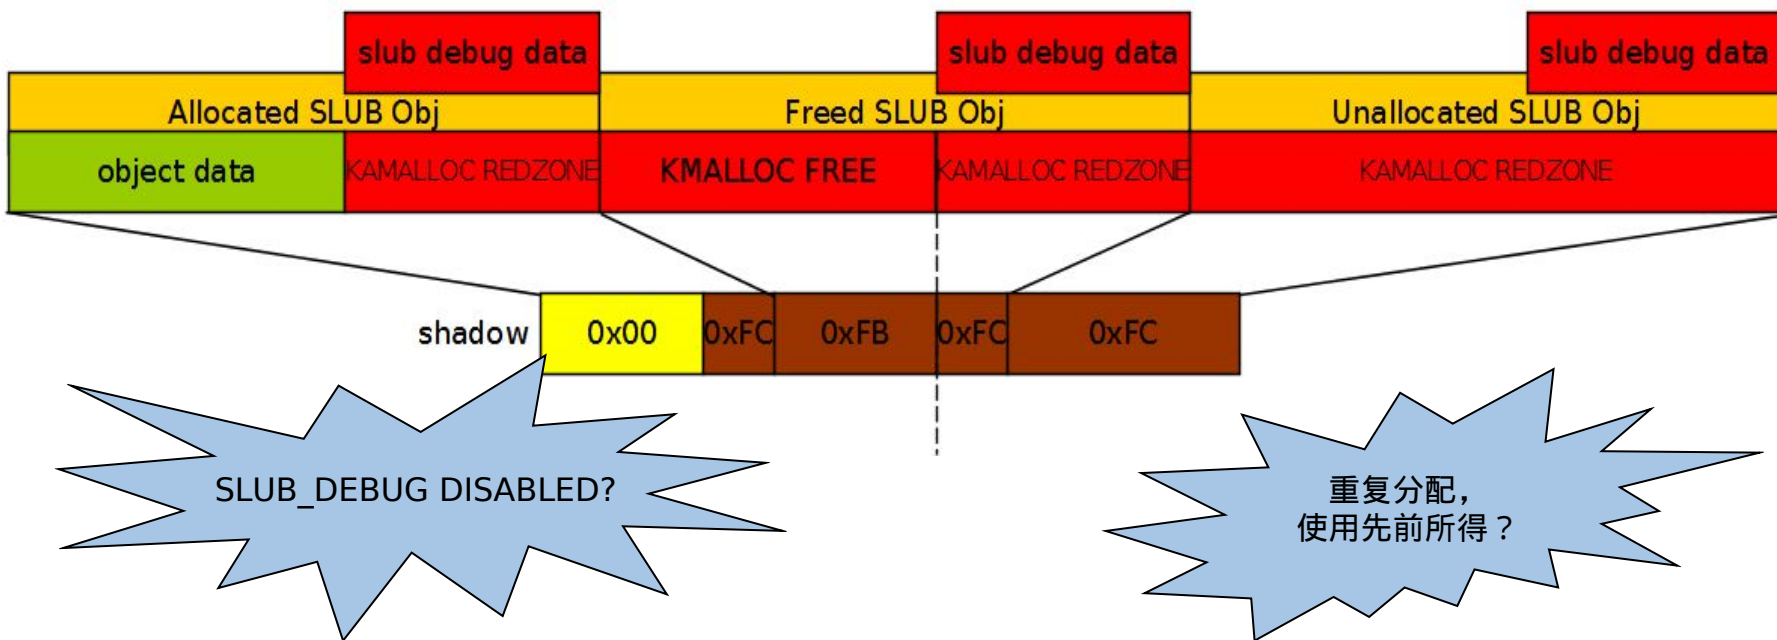


原理及实现

✦ 实现 (ARM 32bits) (续5)

5. 对 SLUB Object 的check支持(续)

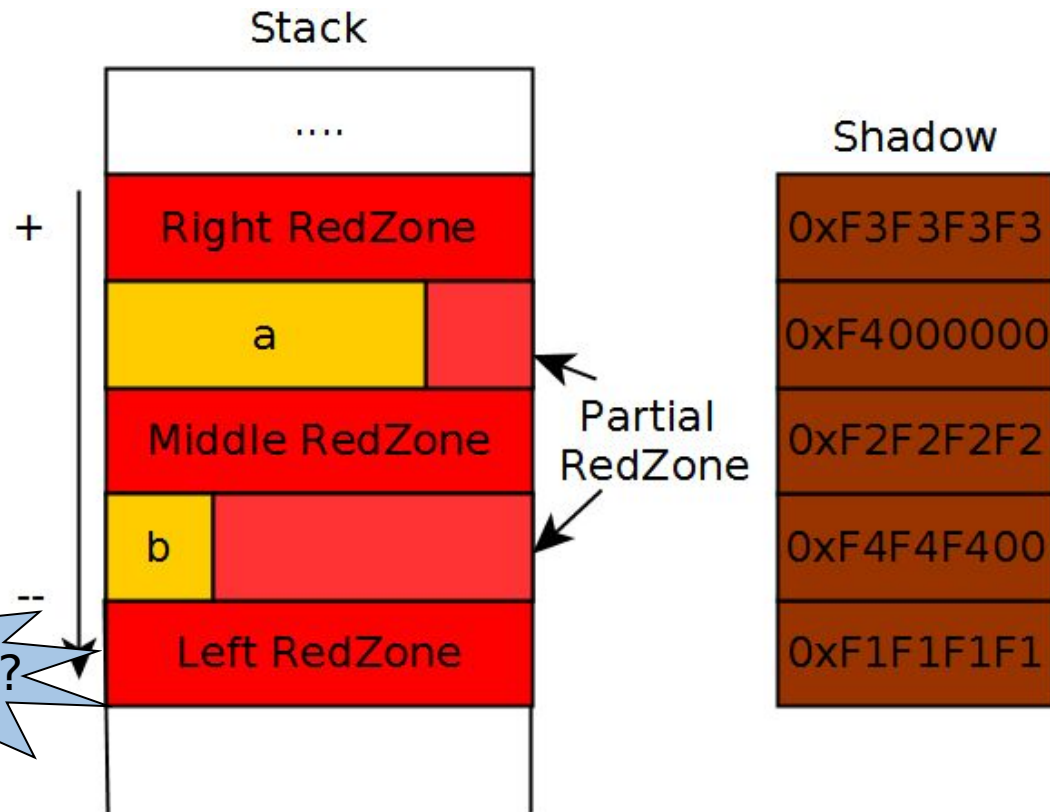

b. 小于等于 8K 的对象



原理及实现

✦ 实现 (ARM 32bits) (续6)
6. 对stack中变量的check支持

```
int  
foo ()  
{  
    char a[23] = {0};  
    int b[2] = {0};  
  
    a[5] = 1;  
    b[1] = 2;  
  
    return a[5] + b[1];  
}
```



原理及实现

✦ 实现 (ARM 32bits) (续7)

7. 对全局变量的check支持(!)

1. 变量要求 32Bytes 对齐;
2. 在变量中间插入 Redzone (0xFA)

```
void __asan_register_globals(struct kasan_global *globals, size_t size)
{
    int i;

    for (i = 0; i < size; i++)
        register_global(&globals[i]);
}
```

```
/* Call all constructor functions linked in */
static void __init do_ctors(void) EXPORT_SYMBOL(__asan_register_globals);
```

```
{
#ifdef CONFIG_CONSTRUCTORS
    ctor_fn_t *fn = (ctor_fn_t *) __ctors_start;

    for (; fn < (ctor_fn_t *) __ctors_end; fn++)
        (*fn)();
#endif
}
```

```
c097ee04 <_GLOBAL__sub_I_65535_1_km_flag>:
c097ee04: e1a0c00d    mov     ip, sp
c097ee08: e92dd800    push   {fp, ip, lr, pc}
c097ee0c: e24cb004    sub     fp, ip, #4
c097ee10: e3a01005    mov     r1, #5
c097ee14: e59f0004    ldr     r0, [pc, #4] ; c097ee20 <_GLOBAL__sub_I_65535_1_km_flag+0x1c>
c097ee18: ebd8f7f5    bl      c0160df4 <__asan_register_globals>
c097ee1c: e89da800    ldm     sp, {fp, sp, pc}
c097ee20: c1071830    .word   0xc1071830
```

```
arm_crash_gcc5> struct kasan_global c1071830 -x
struct kasan_global {
    beg = 0xc0a71040 <__func__.16354>,
    size = 0x11,
    size_with_redzone = 0x40,
    name = 0xc0d811b4,
    module_name = 0xc0da7f83,
    has_dynamic_init = 0x0,
    location = 0xc1071800
}
```

原理及实现

- ✦ 实现 (ARM 32bits) (续8)
- 8. 与其他内存调试工具的比较

Kasan 与其他工具的比较				
	kmemcheck	kasan	SLUB_DEBUG	DEBUG_PAGEALLOC
实现策略	Page Fault Dynamically	Compiler Instrumentation Statically / Redzone	Redzone / Poison / Use-tracing	Compound Page Poison / Unmap
功能	1, Uninit Read 2, UAF 3, OOB	1, UAF 2, OOB	1, UAF 2, OOB 3, User-tracing	1, UAF
速度/Overhead/缺点	最慢 最费内存	较慢 较费内存 不能检查某些 oob	较快 费较少内存 不能检查某些 bad reads 出错点和报告点 在时空上有脱节	较快 不费内存 不能检查细粒度错误

patch 整合

- ✦ 合适的 Toolchain 版本
 1. 最少要求 $\geq 4.9.2$
 2. 若需支持以下功能，则要求 ≥ 5.0
 - a. check OOB of stack/global vars
 - b. inline instrumentation

Sabermud Toolchain / Outline

- ✦ Flash 中 boot 分区大小
 1. boot.img 大小
 - a. $\sim 21\text{M}$ (outline)
 - b. $\sim 32\text{M}$ (inline)
- ✦ U-boot 相关影响？

使用

✦ 配置

1. 要求:

```
CONFIG_KASAN=y  
CONFIG_KASAN_[OUTLINE|INLINE]
```

2. 建议:

```
CONFIG_STACKTRACE=y
```

✦ 命令行

1. 最起码设置 `slub_debug=U`

在 dts 中设置 bootargs

✦ 屏蔽 Kasan

1. 单个文件:

```
KASAN_SANITIZE_main.o := n
```

2. 整个目录:

```
KASAN_SANITIZE := n
```


使用

✧ 例子

```
static noline void __init kmalloc_uaf(void)
{
    char *ptr;
    size_t size = 10;
    unsigned int a,b;

    pr_info("use-after-free\n");
    ptr = kmalloc(size, GFP_KERNEL);
    if (!ptr) {
        pr_err("Allocation failed\n");
        return;
    }

    kfree(ptr);
    *(ptr + 8) = 'x';
}
```

```
c0ed6a38 <kmalloc_uaf>:
c0ed6a38: e1a0c00d    mov     ip, sp
c0ed6a3c: e92dd830    push   {r4, r5, fp, ip, lr, pc}
c0ed6a40: e24cb004    sub    fp, ip, #4
c0ed6a44: e59f1050    ldr     r1, [pc, #80] ; c0ed6a9c <kmalloc_uaf+0x64>
c0ed6a48: e2810040    add    r0, r1, #64 ; 0x40
c0ed6a4c: ebea3db8    bl     c0966134 <printk>
c0ed6a50: e59f0048    ldr     r0, [pc, #72] ; c0ed6aa0 <kmalloc_uaf+0x68>
c0ed6a54: ebca2954    bl     c0160fac <__asan_load4>
c0ed6a58: e59f3044    ldr     r3, [pc, #68] ; c0ed6aa4 <kmalloc_uaf+0x6c>
c0ed6a5c: e3a0200a    mov     r2, #10
c0ed6a60: e3a010d0    mov     r1, #208 ; 0xd0
c0ed6a64: e5930018    ldr     r0, [r3, #24]
c0ed6a68: ebca1dc6    bl     c015e188 <kmem_cache_alloc_trace>
c0ed6a6c: e2504000    subs    r4, r0, #0
c0ed6a70: 1a000003    bne     c0ed6a84 <kmalloc_uaf+0x4c>
c0ed6a74: e59f1020    ldr     r1, [pc, #32] ; c0ed6a9c <kmalloc_uaf+0x64>
c0ed6a78: e28100a0    add    r0, r1, #160 ; 0xa0
c0ed6a7c: ebea3dac    bl     c0966134 <printk>
c0ed6a80: e89da830    ldm     sp, {r4, r5, fp, sp, pc}
c0ed6a84: ebca221e    bl     c015f304 <kfree>
c0ed6a88: e2840008    add     r0, r4, #8
c0ed6a8c: ebca291c    bl     c0160f04 <__asan_store1>
c0ed6a90: e3a03078    mov     r3, #120 ; 0x78
c0ed6a94: e5c43008    strb    r3, [r4, #8]
c0ed6a98: e89da830    ldm     sp, {r4, r5, fp, sp, pc}
c0ed6a9c: c0a71080    .word   0xc0a71080
c0ed6aa0: c12168a0    .word   0xc12168a0
c0ed6aa4: c1216888    .word   0xc1216888
```

```
struct kmem_cache *kmalloc_caches[KMALLOC_SHIFT_HIGH + 1];
EXPORT_SYMBOL(kmalloc_caches);
```



```
[ 14.654296] c0 kasan test: kmalloc_uaf use-after-free
[ 14.654602] c0 =====
[ 14.654693] c0 BUG: KASan: use after free in kmalloc_uaf+0x58/0x70 at addr ecac2808
[ 14.654785] c0 Write of size 1 by task swapper/0/1
[ 14.654846] c0 =====
[ 14.654937] c0 BUG kmalloc-64 (Not tainted): kasan: bad access detected
[ 14.654968] c0 -----
```

```
[ 14.655059] c0 Disabling lock debugging due to kernel taint
[ 14.655181] c0 INFO: Allocated in kmalloc_uaf+0x34/0x70 age=0 cpu=0 pid=1
[ 14.655303] c0      kmem_cache_alloc_trace+0x80/0x22c
[ 14.655395] c0      kmalloc_uaf+0x34/0x70
[ 14.655487] c0      kmalloc_tests_init+0x10/0x34
[ 14.655578] c0      ..... //省略
[ 14.655975] c0      ret_from_fork+0x14/0x20
[ 14.656097] c0 INFO: Freed in kmalloc_uaf+0x50/0x70 age=0 cpu=
[ 14.656188] c0      kfree+0x204/0x214
[ 14.656280] c0      kmalloc_uaf+0x50/0x70
[ 14.656372] c0      kmalloc_tests_init+0x10/0x34
[ 14.656463] c0      ..... //省略
[ 14.656860] c0      ret_from_fork+0x14/0x20
[ 14.656921] c0 INFO: Slab 0xc192e840 objects=16 used=15 fp=0xecac2800 flags=0x0080
[ 14.657012] c0 INFO: Object 0xecac2800 @offset=2048 fp=0x (null)
```

```
[ 14.657135] c0 Bytes b4 ecac27f0: 00 00 00 00 00 00 00 00 00 00 00 00 00 00 00 00 .....
[ 14.657257] c0 Object ecac2800: 00 00 00 00 6e 65 6c 2f 73 6c 61 62 2f 63 66 71 ....nel/slab/cfq
[ 14.657379] c0 Object ecac2810: 5f 69 6f 5f 63 71 00 00 00 00 00 00 00 00 00 00 _io_cq.....
[ 14.657470] c0 Object ecac2820: ..... //省略
[ 14.657806] c0 Padding ecac28f0: 00 00 00 00 00 00 00 00 00 00 00 00 00 00 00 00 .....
[ 14.657928] c0 CPU: 0 PID: 1 Comm: swapper/0 Tainted: G      B      3.10.65-00021-g802eaae-dirty #55
[ 14.658111] c0 [<c001ac34>] (unwind_backtrace+0x0/0x200) from [<c0016d1c>] (show_stack+0x20/0x24)
[ 14.658264] c0 [<c0016d1c>] ..... //省略
[ 14.658874] c0 [<c015cc58>] (object_err+0x3c/0x44) from [<c0161c28>] (kasan_report_error+0x1d4/0x2ac)
[ 14.659057] c0 [<c0161c28>] (kasan_report_error+0x1d4/0x2ac) from [<c0161dfc>] (kasan_report+0x4c/0x54)
[ 14.659210] c0 [<c0161dfc>] (kasan_report+0x4c/0x54) from [<c0160f38>] (__asan_store1+0x34/0x38)
[ 14.659362] c0 [<c0160f38>] (__asan_store1+0x34/0x38) from [<c0ed6a90>] (kmalloc_uaf+0x58/0x70)
[ 14.659515] c0 [<c0ed6a90>] (kmalloc_uaf+0x58/0x70) from [<c0ed6f9c>] (kmalloc_tests_init+0x10/0x34)
[ 14.659698] c0 [<c0ed6f9c>] (kmalloc_tests_init+0x10/0x34) from [<c0009860>] (do_one_initcall_debug+0x80/0x114)
[ 14.659881] c0 [<c0009860>] ..... //省略
[ 14.660430] c0 [<c095d738>] (kernel_init+0x1c/0xf8) from [<c0011388>] (ret_from_fork+0x14/0x20)
[ 14.660491] c0 Memory state around the buggy address:
[ 14.660614] c0  ecac2700: 00 00 00 00 00 00 00 00 fc fc fc fc fc fc fc fc
[ 14.660705] c0  ecac2780: fc fc fc fc fc fc fc fc fc fc fc fc fc fc fc fc
[ 14.660797] c0 >ecac2800: fb fb fb fb fb fb fb fb fc fc fc fc fc fc fc fc
[ 14.660858] c0                ^
[ 14.660949] c0  ecac2880: fc fc fc fc fc fc fc fc fc fc fc fc fc fc fc fc
[ 14.661041] c0  ecac2900: 00 00 00 00 00 00 00 00 fc fc fc fc fc fc fc fc
[ 14.661102] c0 =====
```

```
arm_crash_gcc5> struct kmem_cache ec002140 -x | grep 'size\\|name'
size = 0x100,
object_size = 0x40,
name = 0xec001f00 "kmalloc-64",
name = 0xece47500 "kmalloc-64",
```

参考

✦ kasan 介绍

<https://lwn.net/Articles/612153/>

✦ <https://github.com/torvalds/linux>

✦ <https://github.com/aryabinin/linux>

✦ Porting to GCC 5.0

https://gcc.gnu.org/gcc-5/porting_to.html

✦ <http://www.sabermod.com>